

# Neural Networks - Advanced Image Analysis

---

Jacob Jon Hansen (s134097) Phillip Brinck Vetter (s144097)

## Exercise 1

---




Using python and sklearn we construct our 3 training sets (using our own written functions:

```
#Load data and plot the three cases
plt.figure()
X, y = gen_data("2blop")
plot_data(X,y)

plt.figure()
X, y = gen_data("blop_circle")
plot_data(X,y)

plt.figure()
X, y = gen_data("4blop")
plot_data(X,y)

plt.show()
```



The following code are our high level hyper parameters for all the three networks:

```
#Parameters for NN
num_classes = 2
opt = "adam"
loss_f = 'categorical_crossentropy'
batch_size = 64
nb_epochs = 100
```

Training is run with following code:

```

#Load dataset and fit
X, y = gen_data("2blop")
y_cat = np_utils.to_categorical(y, num_classes)

model = Sequential()
model.add(Dense(3, activation="relu", input_dim=2))
model.add(Dense(5, activation="relu"))
model.add(Dense(2, activation="softmax"))
model.summary()

#Compile the model
model.compile(optimizer = opt, loss = loss_f, metrics=["accuracy"])

hist = model.fit(X, y_cat, batch_size = batch_size,
                 epochs=nb_epochs, verbose=0)

#Plot decision boundary
plot_decision_boundary(X, y, model, cmap='RdBu')



#Plot how the training went
plt.figure()
plot_training(hist)

plt.show()

```



## 2 blops

Layer (type)	Output Shape	Param #
dense_60 (Dense)	(None, 3)	9
dense_61 (Dense)	(None, 5)	20
dense_62 (Dense)	(None, 2)	12
Total params: 41		
Trainable params: 41		
Non-trainable params: 0		





## Blop and circle

Layer (type)	Output Shape	Param #
dense_63 (Dense)	(None, 5)	15
dense_64 (Dense)	(None, 15)	90
dense_65 (Dense)	(None, 10)	160
dense_66 (Dense)	(None, 2)	22
Total params: 287		
Trainable params: 287		
Non-trainable params: 0		



## 4 blops

Layer (type)	Output Shape	Param #
dense_67 (Dense)	(None, 5)	15
dense_68 (Dense)	(None, 15)	90
dense_69 (Dense)	(None, 10)	160
dense_70 (Dense)	(None, 2)	22
Total params: 287		
Trainable params: 287		
Non-trainable params: 0		



## Exercise 2: MNIST

For this we used Keras since we have extensive knowledge of DL and therefore it seemed redundant to sit and program everything from scratch. This is thus not a "valid" competition attempt, but it shows the power of leveraging a good framework to achieve an accuracy of **99.4%**

In order to have a network which generalizes better we use a datagenerator which takes the images from the training set and arguments them with random rotations, zoom and translations:

```
datagen = ImageDataGenerator(rotation_range = 20,  
                             width_shift_range = 0.1,  
                             height_shift_range = 0.1,  
                             zoom_range=0.2)  
  
datagen.fit(X_train)
```

The following model is defined:

Layer (type)	Output Shape	Param #
conv2d_5 (Conv2D)	(None, 26, 26, 32)	320
batch_normalization_7 (Batch Normalization)	(None, 26, 26, 32)	128
activation_7 (Activation)	(None, 26, 26, 32)	0
conv2d_6 (Conv2D)	(None, 24, 24, 32)	9248
batch_normalization_8 (Batch Normalization)	(None, 24, 24, 32)	128
activation_8 (Activation)	(None, 24, 24, 32)	0
max_pooling2d_3 (MaxPooling2D)	(None, 12, 12, 32)	0
conv2d_7 (Conv2D)	(None, 10, 10, 64)	18496
batch_normalization_9 (Batch Normalization)	(None, 10, 10, 64)	256
activation_9 (Activation)	(None, 10, 10, 64)	0
conv2d_8 (Conv2D)	(None, 8, 8, 64)	36928
batch_normalization_10 (Batch Normalization)	(None, 8, 8, 64)	256
activation_10 (Activation)	(None, 8, 8, 64)	0
max_pooling2d_4 (MaxPooling2D)	(None, 4, 4, 64)	0
flatten_2 (Flatten)	(None, 1024)	0
dense_4 (Dense)	(None, 512)	524800
batch_normalization_11 (Batch Normalization)	(None, 512)	2048
activation_11 (Activation)	(None, 512)	0
dropout_3 (Dropout)	(None, 512)	0
dense_5 (Dense)	(None, 256)	131328
batch_normalization_12 (Batch Normalization)	(None, 256)	1024
activation_12 (Activation)	(None, 256)	0
dropout_4 (Dropout)	(None, 256)	0
dense_6 (Dense)	(None, 10)	2570
Total params: 727,530		
Trainable params: 725,610		
Non-trainable params: 1,920		

It uses a CNN, with tricks such as batch normalization (improves training speed) and dropout to increase the ability to generalize to new data.

The network is trained with the following parameters:

```
#Parameters
num_classes = 10
opt = "adam"
loss_f = 'categorical_crossentropy'
batch_size = 64
nb_epochs = 15
```

Training is performed on a big server with a 8 GB GPU, each epoch took 12 seconds (400 on my laptop running a gtx 1050). During training the best model (highest acc on test set) is saved so ensure that we will have the optimal model even if it starts to overfit the training set. This can be done with the following code:

```
#Save best model under training
from keras.callbacks import ModelCheckpoint
filepath="weights_CNN.best.hdf5"
checkpoint = ModelCheckpoint(filepath, monitor='val_acc',
                             verbose=1, save_best_only=True, mode='max')
callbacks_list = [checkpoint]

#Fit the model
model.fit_generator(train_generator,
                    steps_per_epoch = len(X_train) / (2*batch_size),
                    epochs = nb_epochs,
                    validation_data = test_generator,
                    validation_steps = len(y_test)/ batch_size,
                    callbacks = callbacks_list)
```

After training we load the best weights and test:

```
#Evaluate performance of model
model.load_weights("weights_CNN.best.hdf5")
#Compile the model
model.compile(loss=loss_f,
              optimizer=opt,
              metrics=['accuracy'])

score = model.evaluate(X_test, y_test, verbose=0)
print(score[1])

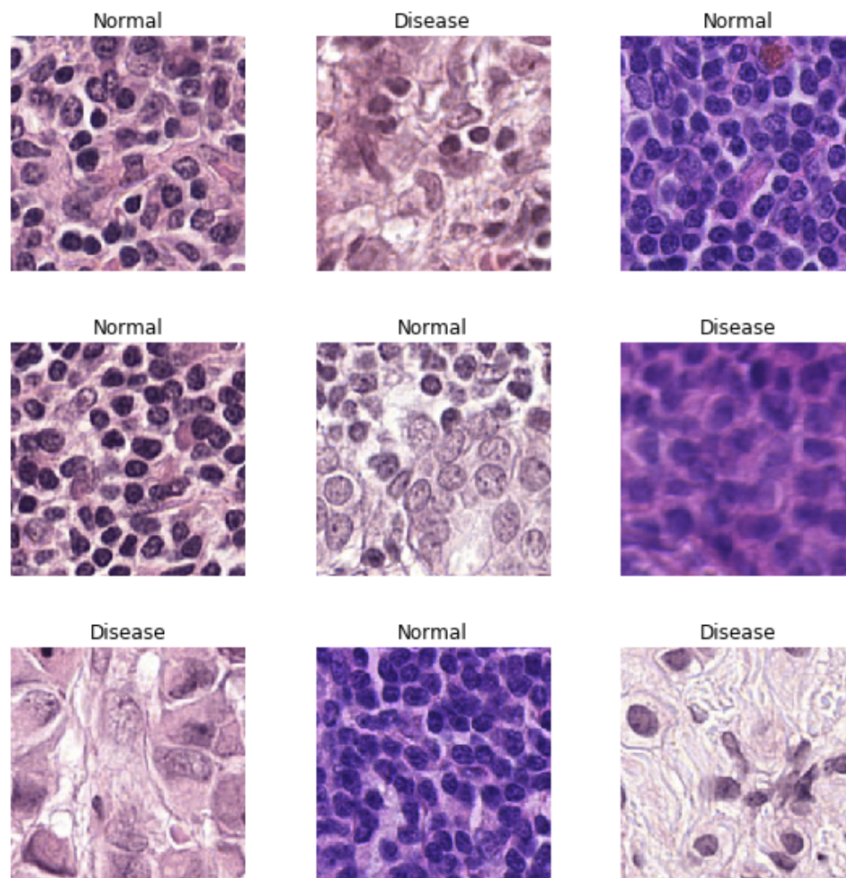
0.9941
```

## Exercise 3: Using VGG19 on breast cancer images.

VGG19 is a powerful architecture that can be downloaded with pretrained weights. The layers in the network are training on the ImageNet challenge, so they will produce a lot of image features. The early layers will give simple features such as lines, circles etc. The last layers will give more abstract features such as eyes, fur, ears etc used for the imagenet dataset. We have used the first fully connected layer at the end and used this to generate features for the images in the breast cancer dataset. A KNN is fitted to these features and we see a performance of **88%** with 11 neighbours. Another method would be to do retraining.

This can be done by freezing all the CNN layers and defining our own output layer and training this with our data to do transfer learning.

### Dataset:



### Model:

Code to load the model (note that we say include\_top=True, this ensure that the model is also using the last FCC layers. We can define a new model from this base model, specifying which layer should be our output.

```
from keras import applications
from keras.models import Sequential, Model

img_width, img_height = (X.shape[0], X.shape[1])

#Define model
model = applications.VGG19(weights = "imagenet", include_top=True,
                             input_shape = (img_width, img_height, 3))
model = Model(inputs=model.input, outputs=model.get_layer('fc1').output)
```

The model has the following architecture, notice that we have 122M parameters so it would take forever to train ourselves:

Layer (type)	Output Shape	Param #
input_9 (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv4 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv4 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv4 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
fc1 (Dense)	(None, 4096)	102764544
Total params: 122,788,928		
Trainable params: 122,788,928		
Non-trainable params: 0		

Now we just predict the features, combine these into a big array and feed them into a KNN.

```

n = X.shape[3]
features_all = np.zeros((n, features.flatten().shape[0]))

#Generate features
for i in range(n):
    x = image.img_to_array(X[:, :, :, i])
    x = np.expand_dims(x, axis=0)
    x = preprocess_input(x)

    features = model.predict(x)
    features_all[i, :] = features.flatten()

from sklearn import neighbors, datasets
from sklearn.metrics import accuracy_score

from sklearn.preprocessing import normalize
normed = normalize(features_all)

knn = neighbors.KNeighborsClassifier(n_neighbors = 11)

knn.fit(normed, y.ravel()-1)

```



```
pred = knn.predict(features_all)
print(accuracy_score(y.ravel()-1, pred)) #-> 0.880
```

